

Finite-Horizon Optimal Control by Differential Dynamic Programming with Application to Robotics

Internship report

Mengda Li

May 16, 2019

Contents

1	Definitions of the problem	3
1.1	Mathematical formulations	3
1.1.1	Transformation of the problem	4
1.1.2	Discretization of a continuous problem	4
1.2	Dynamic Programming	4
2	Differential Dynamic Programming	6
2.1	Decomposition to sub-problems	6
2.2	Backward and Forward pass	7
2.2.1	Forward pass	7
2.2.2	Backward pass	8
2.3	Algorithms	8
2.3.1	Line search methods	9
2.3.2	Regularization methods	9

Introduction

My internship is co-supervised by Justin Carpentier and Nicolas Mansard in the Willow research group at INRIA Paris in France.

Justin Carpentier is a postdoctoral researcher at the interface between Robotics, Machine Learning and Control in Willow. His research is devoted to the embedding of Optimal Control theory inside the formalism of Machine Learning, with Humanoid Robotics as a main target application.

Nicolas Mansard is a permanent researcher in the Gepetto research group at LAAS/CNRS, Toulouse. His research activities are concerned with sensor-based control, and more specifically the integration of sensor-based schemes into humanoid robot applications.

WILLOW is based in the Laboratoire d'Informatique de l'École Normale Supérieure and is a joint research team between INRIA Rocquencourt, École Normale Supérieure de Paris and Centre National de la Recherche Scientifique. Their research is concerned with representational issues in visual object recognition and scene understanding.

GEPETTO is a robotics research team in the Laboratoire d'analyse et d'architecture des systèmes of Centre National de la Recherche Scientifique. They aim at analyzing and generating the motion of anthropomorphic systems.

Chapter 1

Definitions of the problem

We want to find a path from some point x_0 to some point p which doesn't cost too much.

1.1 Mathematical formulations

Let $n \in \mathbb{N}$, $m \in \mathbb{N}$, $x_0 \in \mathbb{R}^n$, $p \in \mathbb{R}^n$, $x \in \mathcal{C}^1([0, T] \rightarrow \mathbb{R}^n)$, $u \in \mathcal{L}^\infty([0, T] \rightarrow \mathbb{R}^m)$, and $f \in \mathcal{C}^2(\mathbb{R}^n \times \mathbb{R}^m \rightarrow \mathbb{R}^n)$.

n is the dimension of the state variable and m is the dimension of the control variable. We use x as the state variable, u as the control variable.

x_0 is the initial position (or state) and p is the desired final position.

x is an unknown function from $[0, T]$ to \mathbb{R}^n . It represents the trajectory.

u is the control function which represents the variation of control over time.

We can control the trajectory x by u :

$$\dot{x}(t) = f(x(t), u(t)) \quad (1.1)$$

the derivative of x depends on x , u and a function f .

Goal 1: Controllability

$$\begin{aligned} &\text{find} && u \\ &\text{subject to} && x(0) = x_0, x(T) = p, \\ &&& \dot{x}(t) = f(x(t), u(t)). \end{aligned} \quad (1.2)$$

Goal 2: Optimal Control

In reality, everything has a cost.

Controlling something has a cost, being somewhere has a cost. So we define a cost (or loss) function l on x and u , supposing that l is in $\mathcal{C}^2(\mathbb{R}^n \times \mathbb{R}^m \rightarrow \mathbb{R})$.

$$\begin{aligned} &\underset{u}{\text{minimize}} && \int_0^T l(x(t), u(t)) dt \\ &\text{subject to} && x(0) = x_0, x(T) = p, \\ &&& \dot{x}(t) = f(x(t), u(t)). \end{aligned} \quad (1.3)$$

1.1.1 Transformation of the problem

Satisfying $x(0) = x_0$ is easy while satisfying $x(T) = p$ is not trivial considering $\dot{x}(t) = f(x(t), u(t))$. To ensure that $x(T) = p$, we can define a final position loss function l_f on x , $l_f \in \mathcal{C}^2(\mathbb{R}^n \rightarrow \mathbb{R})$ and transform the problem to:

$$\begin{aligned} & \underset{u}{\text{minimize}} && \int_{[0, T[} l(x(t), u(t)) dt + l_f(x(T)) \\ & \text{subject to} && x(0) = x_0, \\ & && \dot{x}(t) = f(x(t), u(t)). \end{aligned} \tag{1.4}$$

For example, $l_f(x) = \frac{w}{2} \|x(T) - p^*\|$. w is a weight parameter (like 10^6).

1.1.2 Discretization of a continuous problem

To be computable, we discretize the problem with an horizon $N \in \mathbb{N}^*$, let $h = \frac{T}{N}$, $i \in \{0, \dots, N-1\}$:

$$\begin{aligned} x_{i+1} &= x_i + h\dot{x}_i \\ &= x_i + hf(x_i, u_i) \\ &= F(x_i, u_i) \end{aligned}$$

$$J(U) = \sum_{i=0}^{N-1} L(x_i, u_i) + L_f(x_N) \approx \int_{[0, T[} l(x(t), u(t)) dt + l_f(x(T)) \tag{1.5}$$

where $F(x_i, u_i) = x_i + hf(x_i, u_i)$, $L(x_i, u_i) = hl(x_i, u_i)$, $L_f = l_f$.

Our goal is then to minimize the J over $\ell_N^\infty \subset \mathbb{R}^N$ - a finite-dimensional vector space where the control sequence $U = (u_0, \dots, u_{N-1})$ lies. Here, we transform a problem of infinite-dimensional control to a finite one.

1.2 Dynamic Programming

Notice that:

$$\begin{aligned} \min_U J(U) &= \min_{u_0} \min_{u_1} \dots \min_{u_{N-1}} J(U) \\ &= \min_{u_0} L(x_0, u_0) \\ &\quad + \left(\min_{u_1} L(x_1, u_1) + \dots \min_{u_{N-1}} L(x_{N-1}, u_{N-1}) + L_f(F(x_{N-1}, u_{N-1})) \right) \end{aligned} \tag{1.6}$$

Optimizing directly J over the whole U may be difficult. Following the principle of dynamic programming, we can optimize only over u_i one by one if all future choices u_{i+1}, \dots, u_{N-1} are optimal.

We define the optimal value function V and a function Q such that $V(x) = \min_u Q(x, u)$.

$$\begin{aligned} V_0(x_0) &= \min_{u_0} \min_{u_1} \dots \min_{u_{N-1}} J(U) \\ &= \min_{u_0} L(x_0, u_0) + \underbrace{\min_{u_1} \dots \min_{u_{N-1}} L(x_1, u_1) + \dots + L_f(x_N)}_{V_1(x_1)} \end{aligned} \tag{1.7}$$

$$Q_0(x_0, u_0) = L(x_0, u_0) + \underbrace{\min_{u_1} \dots \min_{u_{N-1}} L(x_1, u_1) + \dots + L_f(x_N)}_{V_1(x_0)} \quad (1.8)$$

For all $i \in \{0, 1, \dots, N-1\}$ iteratively:

$$\begin{aligned} V_i(x_i, u_i) &= \min_{u_i} L(x_i, u_i) + V_{i+1}(x_{i+1}) \\ V_N(x_N) &= L_f(x_N) \end{aligned} \quad (1.9)$$

$$\begin{aligned} Q_i(x_i, u_i) &= L(x_i, u_i) + V_{i+1}(x_{i+1}) \\ &= L(x_i, u_i) + V_{i+1}(f(x_i, u_i)) \\ &= L(x_i, u_i) + \min_{u_{i+1}} Q_{i+1}(f(x_i, u_i), u_{i+1}), \quad i \leq N-2 \end{aligned} \quad (1.10)$$

where

$$V_i(x_i) = \min_{u_i} Q_i(x_i, u_i) \quad (1.11)$$

- V_i is the optimal value of the cost in the sub-trajectory in the time interval $[t_i, T]$ which depends only on the initial position of the sub-trajectory x_i .
- Q_i is the “raw form” of V_i before minimizing over u_i considering that the future controls u_{i+1}, \dots, u_{N-1} are always optimal.

Remark. In practice, what we use in the DDP is the quadratic approximation of the function Q , and that is why we name it Q .

Chapter 2

Differential Dynamic Programming

The computation of differentials (partial derivatives, gradient, and hessian...) is the core of optimization once the function we optimize is differentiable.

Remind that our goal is to optimize $J \in \mathcal{C}^2(\mathbb{R}^N \rightarrow \mathbb{R})$.

$$\begin{aligned} J(u_0, \dots, u_{N-1}) = & L(x_0, u_0) + L(F(x_0, u_0), u_1) + \\ & L(F(F(x_0, u_0), u_1), u_2) + \\ & L_f[F(F(\dots F(F(F(x_0, u_0), u_1), u_2), \dots, u_{N-2}), u_{N-1})] \end{aligned} \quad (2.1)$$

J is differentiable. However, the computation of differentials is not always easy because of the compositions of the “dynamic system” of F . Differential dynamic programming can avoid this hardcore work.

2.1 Decomposition to sub-problems

Remind our final goal is, for all x_0 ,

Goal in the general form

$$\begin{aligned} \text{find} \quad & U^* = (u_0^*, u_1^*, \dots, u_{N-1}^*) \\ \text{subject to} \quad & x(0) = x_0, \\ & J(U^*) = \min_{u_0} \min_{u_1} \dots \min_{u_{N-1}} J(U) \end{aligned} \quad (2.2)$$

which is equivalent to, for all $i \in \{0, 1, \dots, N-1\}$,

Goals in the DDP form

$$\begin{aligned} \text{find} \quad & u_i^* \\ \text{subject to} \quad & x(t_i) = x_i, \\ & Q_i(x_i, u_i^*) = \min_{u_i} Q(x_i, u_i) = V_i(x_i) \end{aligned} \quad (2.3)$$

where $t_0 = 0$, $t_{i+1} = t_i + h$

We will first pre-compute what we need to find each “minimum” of Q_i (their gradients and Hessians) from $N - 1$ to 0 which is called a *Backward Pass*. Once we have these information from differentials, we will compute the “minimum” from 0 to N which is called a *Forward Pass*.

The reason why we first go from future to present is the convenience of computing of differentials to find the minimum. However, we cannot find them immediately in the future: we need to wait the time goes by to compute the minimum from present one by one.

Remark. *We don't search the exact global minimum and we don't even get a local minimum by doing backward and forward pass if the functions Q_i are not quadratic or linear. However, there are some proofs show that by doing backward pass and forward pass iteratively, it will converge to a local minimum of J .*

2.2 Backward and Forward pass

If $Q, V \in \mathcal{C}^2$, then we can proceed the following DDP steps:

2.2.1 Forward pass

For some fixed x_i, u_i

$$V_i(x_i) = \min_u Q_i(x_i, u) = Q_i(x_i, u_i) + \min_{\delta u} [Q_i(x_i, u_i + \delta u) - Q_i(x_i, u_i)]$$

We approximate Q_i quadratically in practice, because *doing global optimization of a function without special properties is hopeless*¹, and approximating the function in higher orders Taylor series is computationally expensive:

$$Q_i(x_i + \delta_x, u_i + \delta_u) - Q_i(x_i, u_i) \approx DQ_i(x_i, u_i)(\delta_x, \delta_u) + \frac{1}{2}D^2Q_i(x_i, u_i)(\delta_x, \delta_u)(\delta_x, \delta_u) \quad (2.4)$$

What we need is then to find the optimal δ_u^* minimizing $Q_i(x_i + \delta_x, u_i + \delta_u)$. **Fixing** x_i, u_i, δ_x and optimizing δ_u , we have

$$\delta_u^*(i) = -(D_{uu}^2 Q_i(x_i, u_i))^{-1}(\nabla_u Q_i(x_i, u_i) + D_{ux}^2 Q_i(x_i, u_i)\delta_x) \quad (2.5)$$

By abuse of notation, we note

$$\bullet \quad \delta_u^* = -Q_{uu}^{-1}(Q_u + Q_{ux}\delta_x) \quad (2.6)$$

with the open-loop term:

$$k = -Q_{uu}^{-1}Q_u \quad (2.7)$$

and the feedback gain term:

$$K = -Q_{uu}^{-1}Q_{ux} \quad (2.8)$$

To compute the best variation of control variable δ_u^* , we need to know the differential of Q :

$$DQ_i(x_i, u_i) = DL(x_i, u_i) + \underline{DV_{i+1}}(x_{i+1}) \circ DF(x, u) \quad (2.9)$$

where the unknown is underlined. However, what we need is only

$$\bullet \quad Q_u = L_u + \underline{V_x'} F_u \quad (2.10)$$

¹by Francis Bach, 2019

¹In the community, we use the prime notation to mean the “next”

And the hessian which can be ignored in iLQR but not in DDP ³

$$\begin{aligned} D^2Q_i(x_i, u_i)(e_k)(e_l) &= D^2L(x_i, u_i)(e_k)(e_l) + \\ &\quad \underline{DV_{i+1}}(x_{i+1})(D^2F(x_i, u_i)(e_k)(e_l)) + \\ &\quad \underline{D^2V_{i+1}}(x_{i+1})(DF(x_i, u_i)(e_k))(DF(x_i, u_i)(e_l)) \end{aligned} \quad (2.11)$$

⁴ where e_k, e_l are vectors of canonical basis. We need use them as canonical variations to compute the Hessian matrix. In a less rigorous form:

$$\bullet \quad Q_{uu} = L_{uu} + \underline{V'_x} \cdot F_{uu} + F_u^T \underline{V'_{xx}} F_u \quad (2.12)$$

$$\bullet \quad Q_{ux} = L_{ux} + \underline{V'_x} \cdot F_{ux} + F_u^T \underline{V'_{xx}} F_x \quad (2.13)$$

where the \cdot denote the tensor dot product. ⁵

As the underline information is from the future, we need the backward pass.

2.2.2 Backward pass

In each backward pass, we compute the first and second order differentials (\approx gradient and hessian) of each V_i and we use this information to find the quadratically approximate minimum of each V_i (thus J).

For all fixed u_i

$$V_i(x_i) = \min_u Q_i(x_i, u) \approx Q_i(x_i, u_i) - \frac{1}{2} D_u Q_i(x_i, u_i) (D_{uu}^2 Q_i(x_i, u_i))^{-1} \nabla_u Q_i(x_i, u_i) \quad (2.14)$$

$$DV_i(x_i) = D_x Q_i(x_i, u_i) - D_u Q_i(x_i, u_i) \circ [D_{uu}^2 Q_i(x_i, u_i)]^{-1} \circ D_{ux}^2 Q_i(x_i, u_i) \quad (2.15)$$

$$D^2V_i(x_i) = D_{xx}^2 Q_i(x_i, u_i) - D_{xu}^2 Q_i(x_i, u_i) \circ [D_{uu}^2 Q_i(x_i, u_i)]^{-1} \circ D_{ux}^2 Q_i(x_i, u_i) \quad (2.16)$$

which are

$$\bullet \quad V_x = Q_x - Q_u Q_{uu}^{-1} Q_{ux} \quad (2.17)$$

$$\bullet \quad V_{xx} = Q_{xx} - Q_{xu} Q_{uu}^{-1} Q_{ux} \quad (2.18)$$

Remark. As we do a quadratic approximation, the higher order (> 2) differentials are supposed to be null. So Q_{uu}^{-1} is constant.

2.3 Algorithms

In a linear or quadratic case (for Q_i), practical results show that by doing the procedure described below only for one time, the trajectory will converge. It is comprehensible because each $\delta_u^*(i)$ is indeed the best step of control variable and we indeed have a real minimum on each $Q_i(x_i, u_i)$. The approximative equality in 2.4 and 2.14 is nor more approximative but exact. The theory guarantees that we have a global minimum of J by doing only one time backward and forward pass. In a non linear and non quadratic case, things are much more complicate: we need iterate this procedure

³[?]

⁴To prove this, use the theorem in the reminder of Differential calculus

⁵It is not a real tensor dot product and F_{uu}, F_{ux} are neither real tensors defined in Wikipedia.

many times until we have a satisfied convergence i.e. the trajectory of control is becoming to satisfy the hypothesis of dynamic programming and the Q_i is becoming “locally quadratic”.

First, we compute the $DV_N(x_N) = DL_{final}(x_N)$ which is directly computable. Once we have $DV_N(x_N)$, we have $DQ_{N-1}(x_{N-1}, u_{N-1})$. Then we can find $DV_{N-1}(x_{N-1})$, and $DQ_{N-2}(x_{N-2}, u_{N-2})$ by the same method. We iterate the process until we have all $DQ_i(x_i, u_i)$ from $N-1$ to 0.

Second, we compute the “best step of control variable” $\delta_u^* = (\delta_u^*(i))_{i \in \{0, \dots, N-1\}}$. If we are in a linear or quadratic case, we update the new control trajectory simply by $u = u + \delta_u$ and then compute the new and converging state trajectory by Euler method. In a non linear and non quadratic case, we need to use line search methods to find a best step in the direction of δ_u^* and regularization methods if the hessian matrices are not “regular”.

2.3.1 Line search methods

2.3.2 Regularization methods

Back to our simple exemple

Initially, we generate a original sequence of x, u : $x_i = 0, u_i = 0 \forall i \in \{1, \dots, N\}$. Restricting in the plan i.e. setting $x, u, p^* \in \mathbb{R}^2$, we have a better graphic representation. Our goal is to make our robot x find a path to go to some place p^* at the end of time T .

In the first iteration, the final on-going loss is clair:

$$V_N^*(x) = V_N(x) = \frac{w}{2} \|x - p^*\|^6$$

$$DV_N^*(x) = w(x^T - p^{*T}) \quad (2.19)$$

$$D^2V_N^*(x) = wId_{2 \times 2} \quad (2.20)$$

In our case,

$$l(x, u) = u^T u$$

$$Dl(x, u) = (0 \quad 0 \quad 2u_1 \quad 2u_2) \quad (2.21)$$

$$D_x l(x, u) = 0, \quad D_u l(x, u) = 2u \quad (2.22)$$

$$D^2 l(x, u) = \begin{pmatrix} 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \quad (2.23)$$

$$D_{uu}^2 l = 2Id_{2 \times 2}, \quad D_{xx}^2 l(x, u) = 0, \quad D_{xu}^2 l(x, u) = 0 \quad (2.24)$$

And

$$Df(x, u) = \begin{pmatrix} 1 & 0 & d & 0 \\ 0 & 1 & 0 & d \end{pmatrix} \quad (2.25)$$

$$D^2 f(x, u) = 0 \quad (2.26)$$

⁶if not mentioned, the norm is euclidean

Reminder of Differential calculus

Theorem 2.3.1 (Second order differential of composition of functions).

$$D^2(f \circ g)(x)(h_1)(h_2) = Df(g(x))(D^2g(x)(h_1)(h_2)) + D^2f(g(x))(Dg(x)(h_1))(Dg(x)(h_2))$$

Proof. By the lemma below. □

For $f : X \rightarrow L(V, W)$, $g : X \rightarrow L(U, V)$

Definition 2.3.1 (generalized product of functions).

$$(f \cdot g)(x) := f(x) \circ g(x)$$

Lemma 2.3.2.

$$D(f \cdot g)(x)(h) = f(x) \circ Dg(x)(h) + Df(x)(h) \circ g(x)$$

Remark. *Books remain silent about this theorem. Either not mentioned or in a false expression.*